

# Class 07: Machine Learning 1

Mitchell Sullivan (PID: A18595276)

## Table of contents

Background . . . . .	1
K-means Clustering . . . . .	2
Hierarchical Clustering . . . . .	7
PCA of UK food data . . . . .	8
Spotting major differences and trends . . . . .	10
Pairs plots and heatmaps . . . . .	12
PCA to the rescue . . . . .	13

## Background

Today we will begin our exploration of some important machine learning methods, namely **clustering** and **dimensionality reduction**.


Let's make up some input data for clustering where we know what the natural clusters are.

The function `rnorm()` can be useful here.

```
rnorm3pos <- rnorm(30, mean = 3)
rnorm3neg <- rnorm(30, mean = -3)
```

```
tmp <- c(rnorm3pos, rnorm3neg)
x <- cbind(x = tmp, y = rev(tmp))
```

```
plot(x)
```



## K-means Clustering

The main function in “base R” for K-means clustering is called `kmeans()`

```
km <- kmeans(x = x, centers = 2)  
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.811689	2.792943
2	2.792943	-2.811689

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 67.92911 67.92911
```

(between SS / total SS = 87.4 %)

## Available components:

```
[1] "cluster"        "centers"        "totss"          "withinss"        "tot.withinss"  
[6] "betweenss"      "size"           "iter"           "ifault"
```

Q. What component of the results object details the cluster sizes?

km\$size

[1] 30 30

Q. What component of the results object details the cluster centers?

km\$centers

```

          x         y
1 -2.811689  2.792943
2  2.792943 -2.811689


```

Q. What component of the results object details the cluster membership vector (i.e. our main result of which points lie in which cluster)?

km\$cluster

Q. Plot our sclustering results with points colored by cluster and also add the cluster centers as new points colored blue

```
plot(x, col = km$cluster)
points(km$centers, col = "blue", pch = 15)
```



Q. Run `kmeans()` again and this time produce 4 clusters (call your result object `k4`) and make a results figure like above?

```
k4 <- kmeans(x, centers = 4)
k4
```

K-means clustering with 4 clusters of sizes 13, 18, 17, 12

Cluster means:

	x	y
1	-2.613059	3.868344
2	2.198523	-2.294220
3	-2.963583	1.970578
4	3.684573	-3.587893

Clustering vector:

```
[1] 2 4 2 4 4 4 2 2 4 2 2 2 2 4 4 4 4 2 2 2 2 4 2 4 2 2 3 1 3 3 3 3 3 3  
[39] 3 3 3 1 1 3 1 1 3 3 1 3 3 1 1 1 3 1 1
```


Within cluster sum of squares by cluster:

```
[1] 20.08473 24.23929 20.40808 15.73988
(between_SS / total_SS =  92.5 %)
```

```
Available components:
```

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(x, col = k4$cluster)  
points(k4$centers, col = "blue", pch = 15)
```



The metric total within sums of squares gives us more information on the “fit” of the clustering

```
km$tot.withinss
```

```
[1] 135.8582
```

```
k4$tot.withinss
```

```
[1] 80.47198
```

**Key-point** Because K-means will impose the number of groups you specify, it can artificially form clusters that aren’t there.


Q. Let's try different numbers of K (centers) from 1 to 30 and see what the best result is

```
ans <- NULL
for (i in 1:30) {
  ans <- c(ans, kmeans(x, centers = i)$tot.withinss)
}

ans
```

```
[1] 1078.215265 135.858221 107.908277 80.471975 83.589202 55.639257
[7] 52.712983 35.665568 46.539064 32.624309 20.145512 22.143162
[13] 16.384573 15.277305 14.298740 11.647889 11.524407 9.814527
[19] 9.610789 8.261072 8.303838 5.335909 4.038065 4.998548
[25] 5.272589 3.454351 3.353336 2.889346 2.891236 2.737496
```

```
plot(ans, typ = "o")
```




Using this method, we can find where the biggest step down in total within sums of squares and pick the best number of clusters.

## Hierarchical Clustering


The main function for Hierarchical Clustering is called `hclust()`. Unlike `kmeans()` (which does all the work for you), you can't just pass `hclust()` our raw input data. It needs a "distance matrix" like the one returned from the `dist()` function.

```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```



To extract our cluster membership vector from an `hclust()` result object, we have to "cut" our tree at a given height to yield separate "groups" or "branches".

```
plot(hc)
abline(h = 8, col = "red", lty = 2)
```



To do this, we use the `cutree()` function on our `hclust()` object:

```
grps <- cutree(hc, h = 6)  
grps
```

```
table(grps, km$cluster)
```

grps	1	2
1	0	30
2	30	0

## PCA of UK food data

Import the dataset of food consumption in the UK

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)  
x
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139
7	Fresh_potatoes		720	874	566	1033
8	Fresh_Veg		253	265	171	143
9	Other_Veg		488	570	418	355
10	Processed_potatoes		198	203	220	187
11	Processed_Veg		360	365	337	334
12	Fresh_fruit		1102	1137	957	674
13	Cereals		1472	1582	1462	1494
14	Beverages		57	73	53	47
15	Soft_drinks		1374	1256	1572	1506
16	Alcoholic_drinks		375	475	458	135
17	Confectionery		54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

There are 17 rows and 5 columns. Using the `dim()` function will tell us this.

```
dim(x)
```

```
[1] 17 5
```

One solution to set the row names is to do it by hand...

```
rownames(x) <- x[,1]
head(x)
```

		X	England	Wales	Scotland	N.Ireland
Cheese	Cheese		105	103	103	66
Carcass_meat	Carcass_meat		245	227	242	267
Other_meat	Other_meat		685	803	750	586
Fish	Fish		147	160	122	93
Fats_and_oils	Fats_and_oils		193	235	184	209
Sugars	Sugars		156	175	147	139

```
dim(x)
```

```
[1] 17 5
```

A better way to do this is to set the row names to the first column with `read.csv()`.

```
x <- read.csv(url, row.names = 1)
x
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer using the `read.csv()` argument because it won’t delete a column when run more than once.


## Spotting major differences and trends

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?


```
barplot(as.matrix(x), col=rainbow(nrow(x)))
```




Removing `beside = T` stacks the bars

### Pairs plots and heatmaps

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



## PCA to the rescue

The main PCA function in “base R” is called `prcomp()`. This function wants the transpose of our food data as input (i.e. the foods as columns and the countries as rows).

```
pca <- prcomp(t(x))
```

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"

$class
[1] "prcomp"
```


To make one of our main PCA result figures, we turn to `pca$x` (the scores along our PCs). This is called “PC plot”, “score plot”, or “ordination plot” ...

```
library(ggplot2)
```

```
Warning: package 'ggplot2' was built under R version 4.4.3
```

```
my_cols <- c("orange", "red", "blue", "darkgreen")
```

```
ggplot(pca$x) +
  aes(x = PC1, y = PC2, label = rownames(pca$x), colour = rownames(pca$x)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_minimal() +
  theme(legend.position = "none")
```



The second major result figure is called a “loadings plot”, “variable contributions plot”, or “weight plot”.

```

ggplot(pca$rotation) +
  aes(x = PC1,
      y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "grey30") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_minimal() +
  theme(axis.text.y = element_text(size = 8.5))

```

