

Class09 - Candy Mini-Project

Mitchell Sullivan (PID: A18595276)

Table of contents

| | |
|---|----|
| Background | 1 |
| Data Import | 1 |
| What is in the dataset? | 2 |
| What is your favorite candy? | 3 |
| Using the <code>skimr::skim()</code> function | 4 |
| Exploratory analysis | 5 |
| Overall candy rankings | 7 |
| Adding color | 13 |
| Looking at pricepercent | 15 |
| Exploring the correlation structure | 16 |
| Principal Component Analysis | 17 |
| Summary | 20 |

Background

In today's mini-project, we will analyze candy data with the exploratory graphics, basic statistics, correlation analysis, and principal component analysis methods we have been learning thus far.

Data Import

The data comes as a CSV file from 538.

```
candy <- read.csv("candy-data.csv", row.names = 1)

head(candy)
```

| | chocolate | fruity | caramel | peanuty | almondy | nougat | crisped | rice | wafer | |
|--------------|-----------|----------|--------------|--------------|------------|----------|---------|------|-------|---|
| 100 Grand | 1 | 0 | 1 | | 0 | 0 | | | | 1 |
| 3 Musketeers | 1 | 0 | 0 | | 0 | 1 | | | | 0 |
| One dime | 0 | 0 | 0 | | 0 | 0 | | | | 0 |
| One quarter | 0 | 0 | 0 | | 0 | 0 | | | | 0 |
| Air Heads | 0 | 1 | 0 | | 0 | 0 | | | | 0 |
| Almond Joy | 1 | 0 | 0 | | 1 | 0 | | | | 0 |
| hard | bar | pluribus | sugarpercent | pricepercent | winpercent | | | | | |
| 100 Grand | 0 | 1 | 0 | 0.732 | 0.860 | 66.97173 | | | | |
| 3 Musketeers | 0 | 1 | 0 | 0.604 | 0.511 | 67.60294 | | | | |
| One dime | 0 | 0 | 0 | 0.011 | 0.116 | 32.26109 | | | | |
| One quarter | 0 | 0 | 0 | 0.011 | 0.511 | 46.11650 | | | | |
| Air Heads | 0 | 0 | 0 | 0.906 | 0.511 | 52.34146 | | | | |
| Almond Joy | 0 | 1 | 0 | 0.465 | 0.767 | 50.34755 | | | | |

What is in the dataset?

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

[1] 85

There are 85 rows in this dataset.

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

[1] 38

There are 38 fruity candy types in this dataset

What is your favorite candy?

Q3. What is your favorite candy (other than Twix) in the dataset and what is its `winpercent` value?

My favorite candy in this dataset is "Caramel Apple Pops". Let's see what its win percent is:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from 'package:stats':
```

```
filter, lag
```

```
The following objects are masked from 'package:base':
```

```
intersect, setdiff, setequal, union
```

```
winpercent <- function(data, candytype) {  
  data |>  
    filter(row.names(candy) == candytype) |>  
    select(winpercent)  
}  
  
winpercent(candy, "Caramel Apple Pops")
```

```
winpercent  
Caramel Apple Pops 34.51768
```

The win percentage is 34.5%.

Q4. What is the `winpercent` value for "Kit Kat"?

```
winpercent(candy, "Kit Kat")
```

```
winpercent  
Kit Kat 76.7686
```

The win percentage is 76.8%.

Q5. What is the `winpercent` value for “Tootsie Roll Snack Bars”?

```
winpercent(candy, "Tootsie Roll Snack Bars")
```

```
      winpercent  
Tootsie Roll Snack Bars    49.6535
```

Using the `skimr::skim()` function

`skim()` from the `skimr` package can give a quick overview of a dataset

```
library("skimr")
```

```
Warning: package 'skimr' was built under R version 4.4.3
```

```
skim(candy)
```

Table 1: Data summary

| | |
|------------------------|-------|
| Name | candy |
| Number of rows | 85 |
| Number of columns | 12 |
| Column type frequency: | |
| numeric | 12 |
| Group variables | |
| | None |

Variable type: numeric

| skim_variable | n_missing | complete_rate | mean | sd | p0 | p25 | p50 | p75 | p100 | hist |
|----------------|-----------|---------------|------|------|------|------|------|------|------|------|
| chocolate | 0 | 1 | 0.44 | 0.50 | 0.00 | 0.00 | 0.00 | 1.00 | 1.00 | |
| fruity | 0 | 1 | 0.45 | 0.50 | 0.00 | 0.00 | 0.00 | 1.00 | 1.00 | |
| caramel | 0 | 1 | 0.16 | 0.37 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 | |
| peanutyalymond | 0 | 1 | 0.16 | 0.37 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 | |

| skim_vari- able | n_miss- ing | com- plete_rate | mean | sd | p0 | p25 | p50 | p75 | p100 | hist |
|-----------------------|----------------|--------------------|-------|-------|-------|-------|-------|-------|-------|------|
| nougat | 0 | 1 | 0.08 | 0.28 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 | |
| crispedrice- wafer | 0 | 1 | 0.08 | 0.28 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 | |
| hard | 0 | 1 | 0.18 | 0.38 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 | |
| bar | 0 | 1 | 0.25 | 0.43 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 | |
| pluribus | 0 | 1 | 0.52 | 0.50 | 0.00 | 0.00 | 1.00 | 1.00 | 1.00 | |
| sugarpercent | 0 | 1 | 0.48 | 0.28 | 0.01 | 0.22 | 0.47 | 0.73 | 0.99 | |
| pricepercent | 0 | 1 | 0.47 | 0.29 | 0.01 | 0.26 | 0.47 | 0.65 | 0.98 | |
| winpercent | 0 | 1 | 50.32 | 14.71 | 22.45 | 39.14 | 47.83 | 59.86 | 84.18 | |

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The `winpercent` column looks the most different because it is what looks like a curve while the others are either binary or uniform distributions.

Q7. What do you think a zero and one represent for the `candy$chocolate` column?

The zero represents that a candy does not contain chocolate while a 1 represents that the candy does have chocolate.

Exploratory analysis

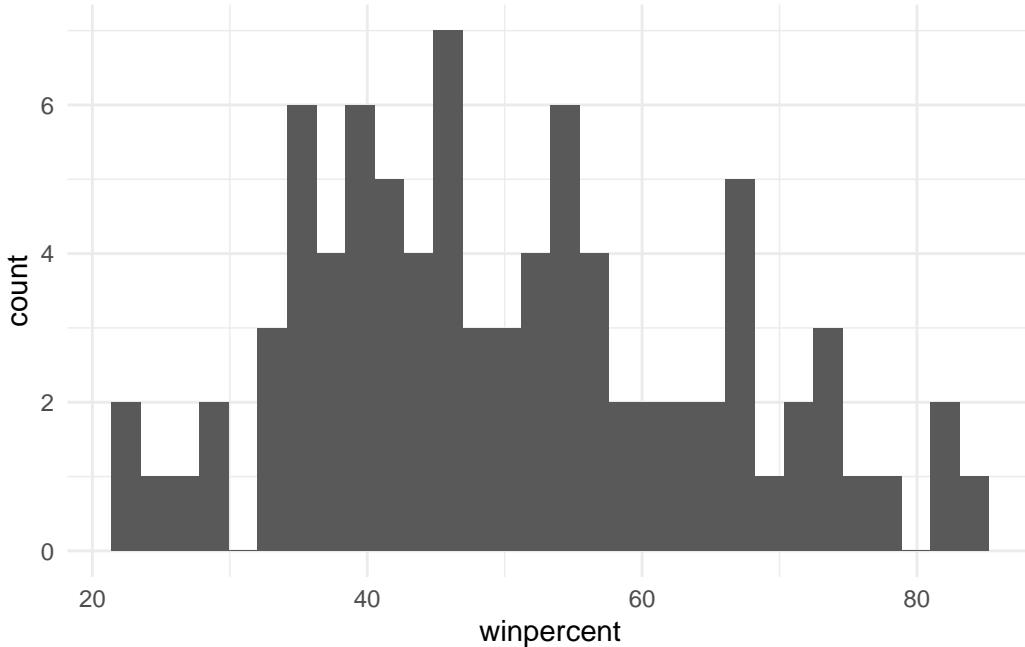
Q8. Plot a histogram of winpercent values using both base R an ggplot2.

```
library(ggplot2)
```

```
Warning: package 'ggplot2' was built under R version 4.4.3
```

```
ggplot(candy) +
  aes(x = winpercent) +
  geom_histogram() +
  theme_minimal()
```

```
`stat_bin()` using `bins = 30`. Pick better value `binwidth`.
```



Q9. Is the distribution of winpercent values symmetrical?

```
mean(candy$winpercent)
```

[1] 50.31676

```
median(candy$winpercent)
```

[1] 47.82975

No, the distribution is skewed right because the median is less than the mean

Q10. Is the center of the distribution above or below 50%?

The center is below 50%

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Steps: 1. Find all chocolate candy in the dataset 2. Extract their `winpercent` values 3. Calculate the mean of these values 4. Repeat for fruity candy in the dataset

```
choc_candy <- candy[candy$chocolate == 1, ]
choc_wins <- choc_candy$winpercent
mean_choc_wins <- mean(choc_wins)
mean_choc_wins
```

```
[1] 60.92153
```

```
fruity_candy <- candy[candy$fruity == 1, ]
fruity_wins <- fruity_candy$winpercent
mean_fruity_wins <- mean(fruity_wins)
mean_fruity_wins
```

```
[1] 44.11974
```

Q12. Is this difference statistically significant?

```
t.test(choc_wins, fruity_wins)
```

```
Welch Two Sample t-test

data: choc_wins and fruity_wins
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Yes, it is statistically significant with a Welch Two Sample t-test.

Overall candy rankings

Q13. What are the five least liked candy types in this set?

```
candy |>
  arrange(winpercent) |>
  head(5)
```

| | chocolate | fruity | caramel | peanuty | almondy | nougat |
|--------------------|-----------|---------|---------|---------|---------|----------|
| Nik L Nip | 0 | 1 | 0 | 0 | 0 | 0 |
| Boston Baked Beans | 0 | 0 | 0 | 1 | 0 | 0 |
| Chiclets | 0 | 1 | 0 | 0 | 0 | 0 |
| Super Bubble | 0 | 1 | 0 | 0 | 0 | 0 |
| Jawbusters | 0 | 1 | 0 | 0 | 0 | 0 |
| | crisped | rice | wafer | hard | bar | pluribus |
| | sugar | percent | price | percent | win | percent |
| Nik L Nip | 0 | 0 | 0 | 1 | 0.197 | 0.976 |
| Boston Baked Beans | 0 | 0 | 0 | 1 | 0.313 | 0.511 |
| Chiclets | 0 | 0 | 0 | 1 | 0.046 | 0.325 |
| Super Bubble | 0 | 0 | 0 | 0 | 0.162 | 0.116 |
| Jawbusters | 0 | 1 | 0 | 1 | 0.093 | 0.511 |
| | win | percent | | | | |
| Nik L Nip | 22.44534 | | | | | |
| Boston Baked Beans | 23.41782 | | | | | |
| Chiclets | 24.52499 | | | | | |
| Super Bubble | 27.30386 | | | | | |
| Jawbusters | 28.12744 | | | | | |

The five least liked candies are Nik L Nip, Chiclets, Super Bubble, and Jawbusters

Q14. What are the top 5 all time favorite candy types out of this set?

```
candy |>
  arrange(desc(winpercent)) |>
  head(5)
```

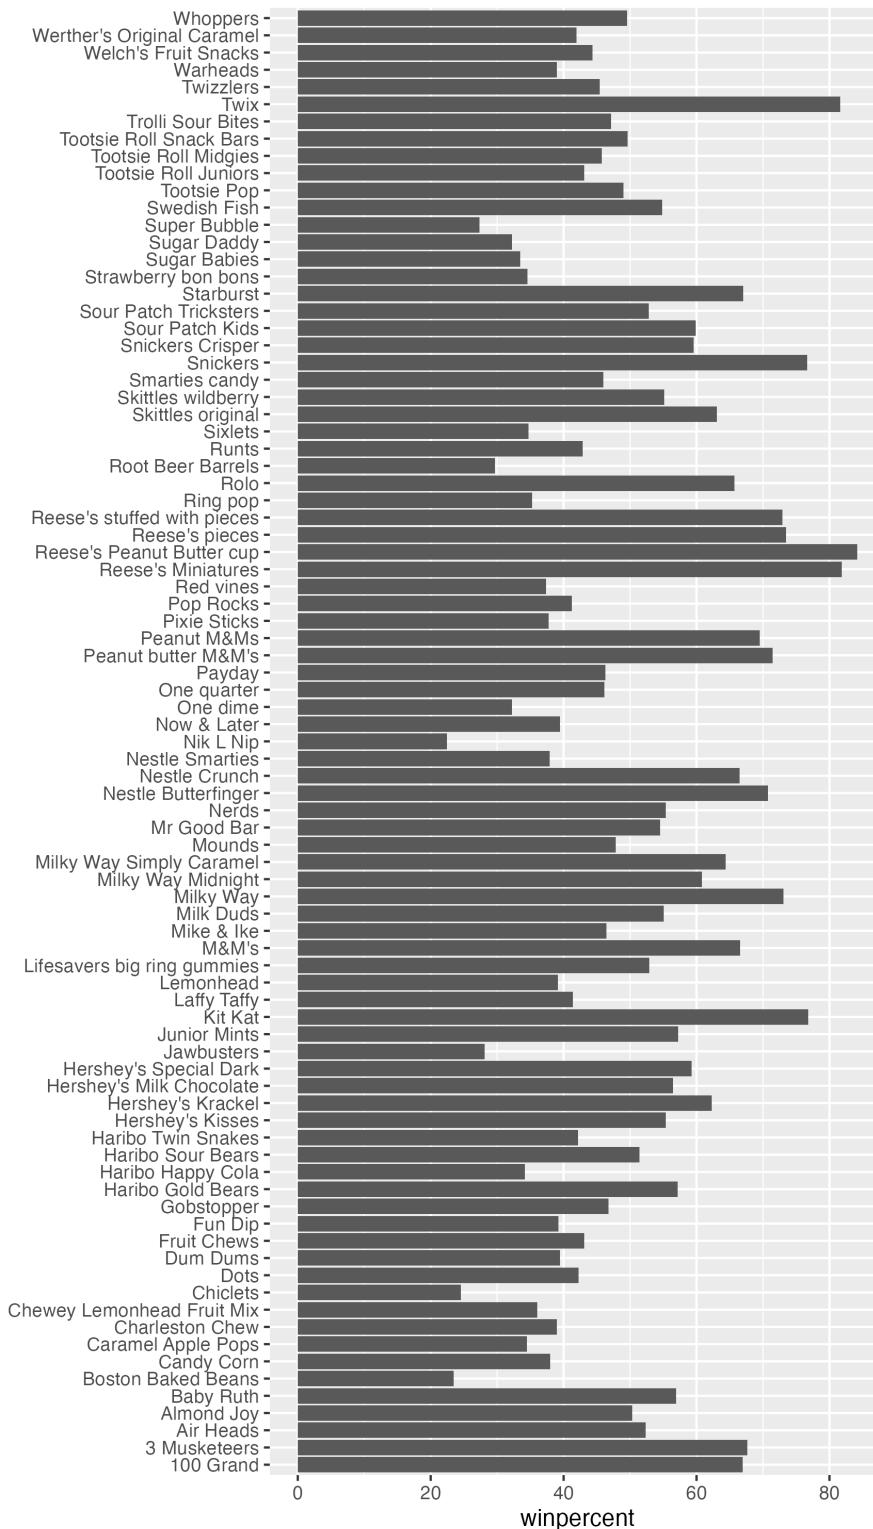
| | chocolate | fruity | caramel | peanuty | almondy | nougat |
|---------------------------|-----------|----------|---------|---------|---------|----------|
| Reese's Peanut Butter cup | 1 | 0 | 0 | 1 | 0 | 0 |
| Reese's Miniatures | 1 | 0 | 0 | 1 | 0 | 0 |
| Twix | 1 | 0 | 1 | 0 | 0 | 0 |
| Kit Kat | 1 | 0 | 0 | 0 | 0 | 0 |
| Snickers | 1 | 0 | 1 | 1 | 1 | 1 |
| | crisped | rice | wafer | hard | bar | pluribus |
| | sugar | percent | price | percent | win | percent |
| Reese's Peanut Butter cup | 0 | 0 | 0 | 0 | 0.720 | 0.720 |
| Reese's Miniatures | 0 | 0 | 0 | 0 | 0.034 | 0.034 |
| Twix | 1 | 0 | 1 | 0 | 0.546 | 0.546 |
| Kit Kat | 1 | 0 | 1 | 0 | 0.313 | 0.313 |
| Snickers | 0 | 0 | 1 | 0 | 0.546 | 0.546 |
| | price | percent | win | percent | | |
| Reese's Peanut Butter cup | 0.651 | 84.18029 | | | | |
| Reese's Miniatures | 0.279 | 81.86626 | | | | |

| | | |
|----------|-------|----------|
| Twix | 0.906 | 81.64291 |
| Kit Kat | 0.511 | 76.76860 |
| Snickers | 0.651 | 76.67378 |

I prefer the `dplyr` version because it is easier to read.

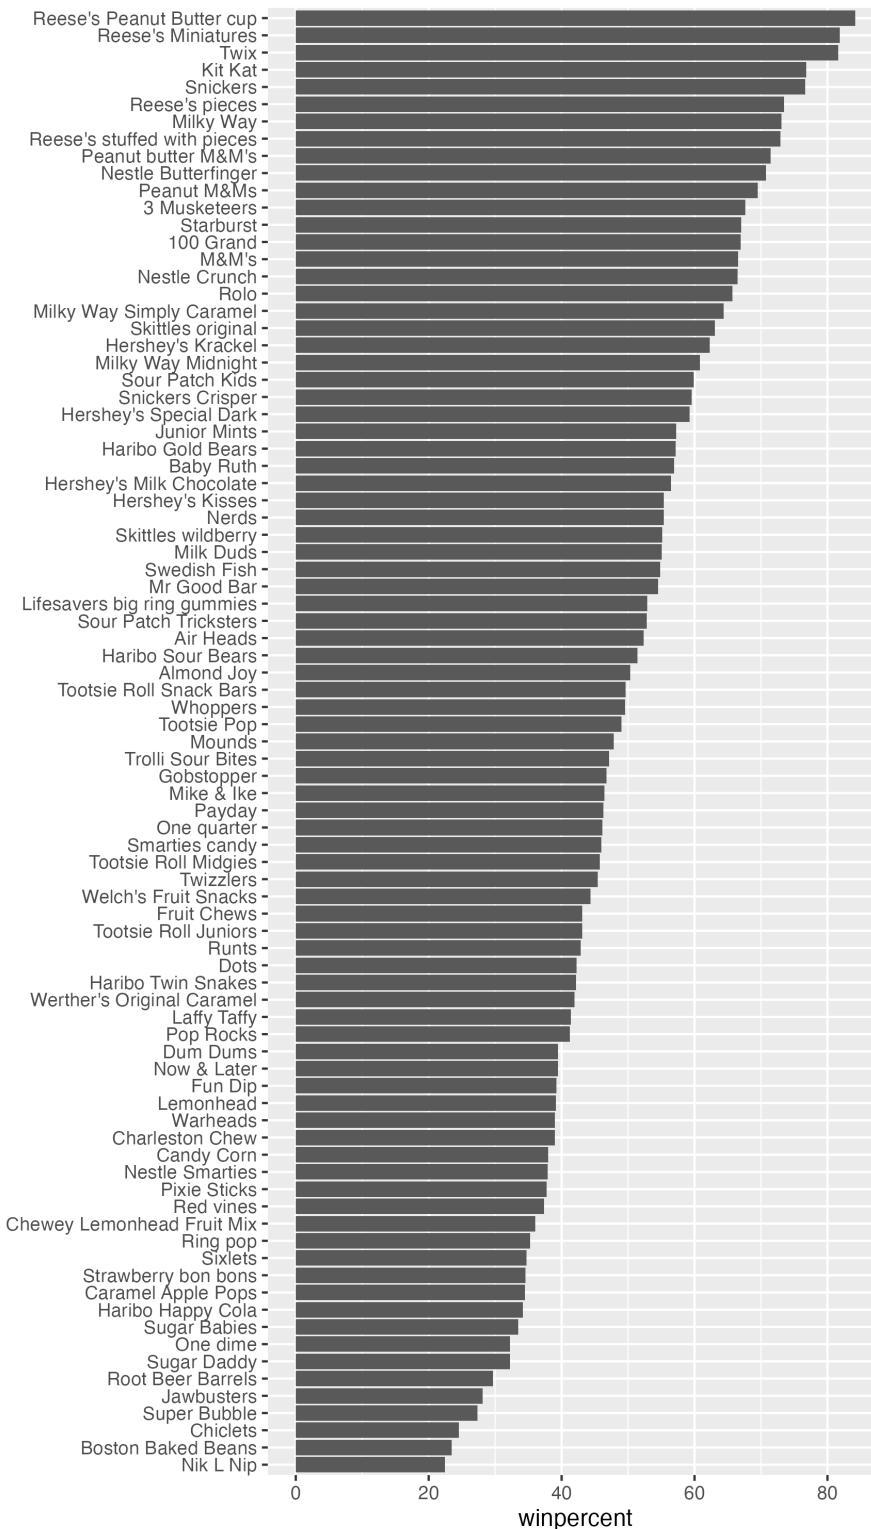
Q15. Make a first barplot of candy ranking based on winpercent values.

```
plot1 <- ggplot(candy) +  
  aes(winpercent, rownames(candy)) +  
  geom_col() +  
  ylab("")  
  
ggsave("barplot1.png", plot = plot1, height = 10, width = 6)
```



Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by `winpercent`?

```
plot2 <- ggplot(candy) +  
  aes(winpercent, reorder(rownames(candy), winpercent)) +  
  geom_col() +  
  ylab("")  
  
ggsave("barplot2.png", plot = plot2, height = 10, width = 6)
```

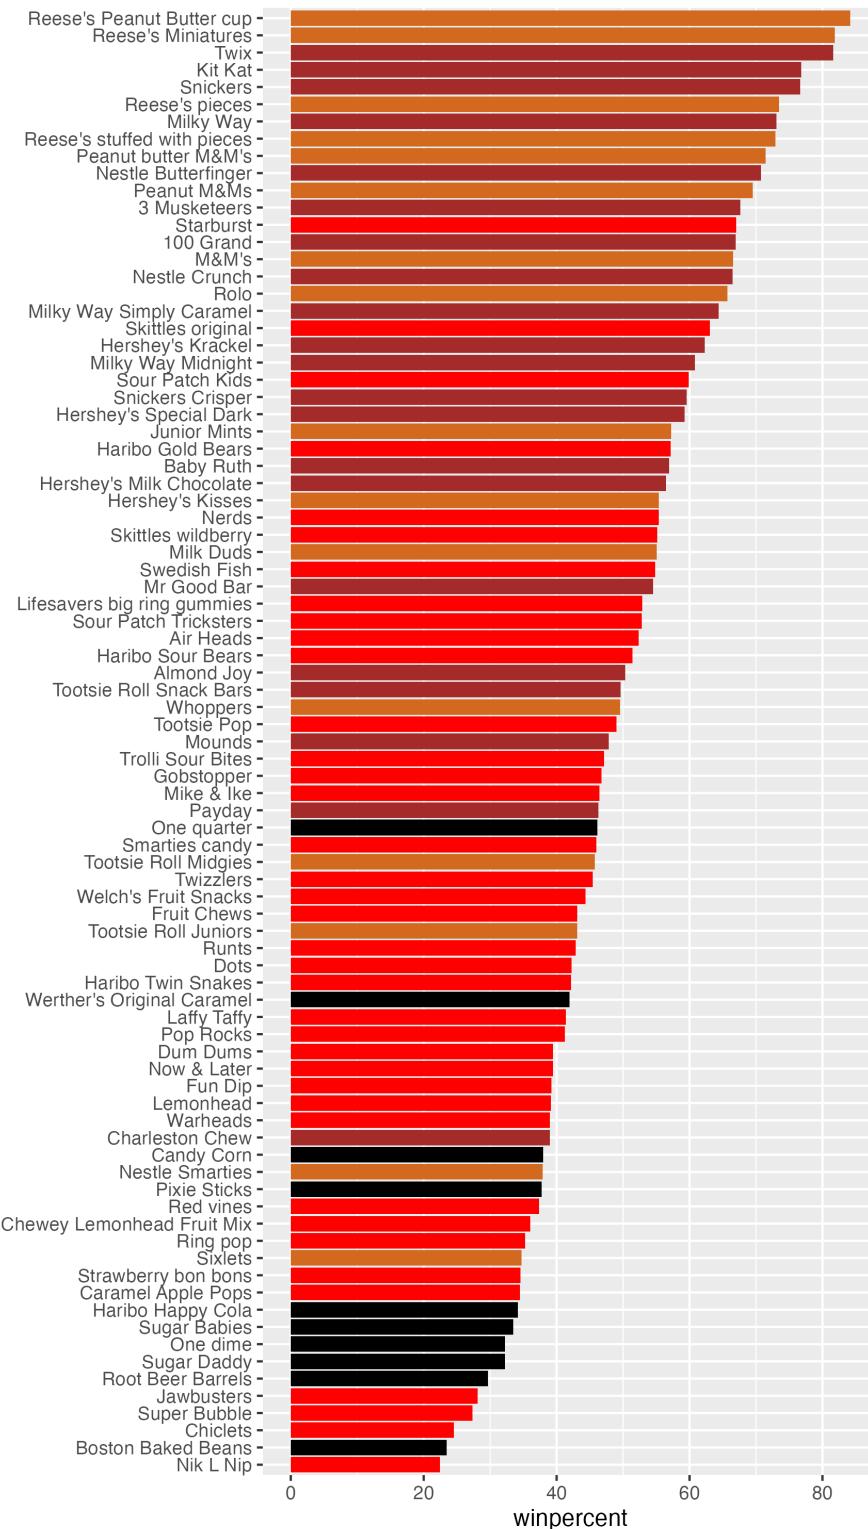


Adding color

```
my_cols <- rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] <- "chocolate"
my_cols[as.logical(candy$bar)] <- "brown"
my_cols[as.logical(candy$fruity)] <- "red"

plot3 <- ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill=my_cols) +
  ylab("")

ggsave("barplot3.png", plot = plot3, height = 10, width = 6)
```



Q17. What is the worst ranked chocolate candy?

Sixlets

Q18. What is the best ranked fruity candy?

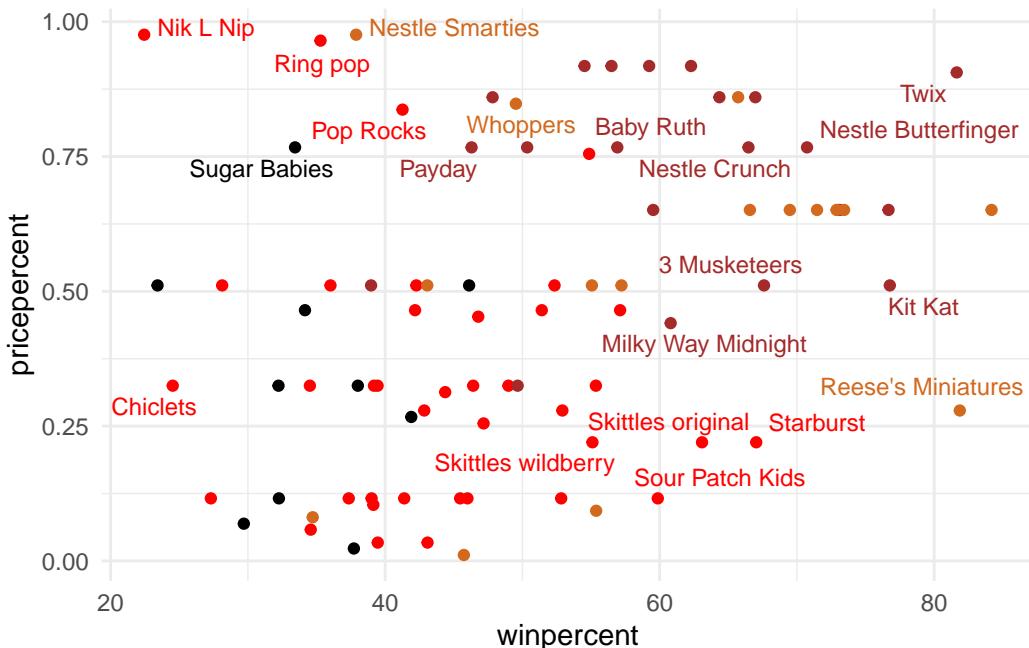
Starburst

Looking at pricepercent

```
library(ggrepel)

# How about a plot of win vs price
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col = my_cols) +
  geom_text_repel(col = my_cols, size = 3.3, max.overlaps = 5) +
  theme_minimal()
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
best_candy <- candy |>
  select(pricepercent, winpercent) |>
  arrange(desc(winpercent))

head(best_candy, 5)
```

| | pricepercent | winpercent |
|---------------------------|--------------|------------|
| Reese's Peanut Butter cup | 0.651 | 84.18029 |
| Reese's Miniatures | 0.279 | 81.86626 |
| Twix | 0.906 | 81.64291 |
| Kit Kat | 0.511 | 76.76860 |
| Snickers | 0.651 | 76.67378 |

Reese's Miniatures is the cheapest of the top 5 winning candies.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
exp_candy <- candy |>
  select(pricepercent, winpercent) |>
  arrange(desc(pricepercent))

head(exp_candy, 5)
```

| | pricepercent | winpercent |
|--------------------------|--------------|------------|
| Nik L Nip | 0.976 | 22.44534 |
| Nestle Smarties | 0.976 | 37.88719 |
| Ring pop | 0.965 | 35.29076 |
| Hershey's Krackel | 0.918 | 62.28448 |
| Hershey's Milk Chocolate | 0.918 | 56.49050 |

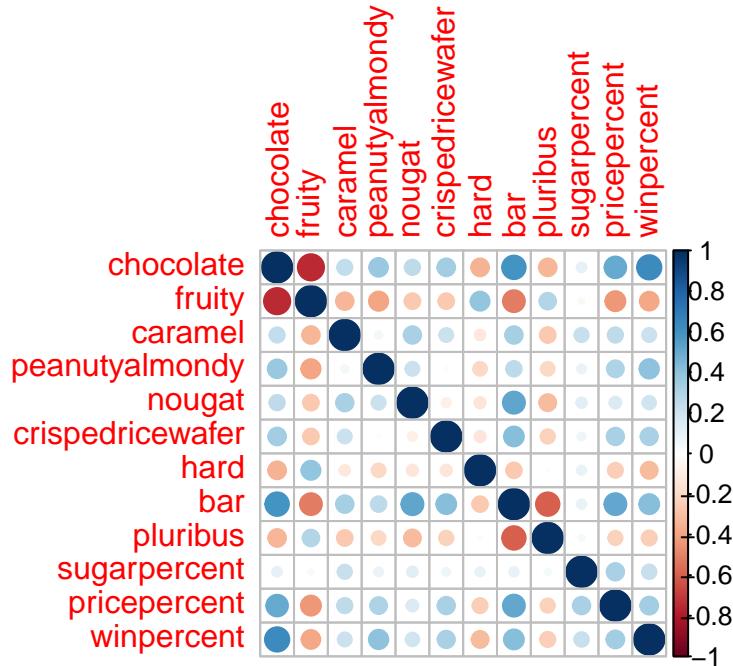
Nik L Nip is least popular of the five most expensive candies.

Exploring the correlation structure

```
library(corrplot)
```

corrplot 0.95 loaded

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and fruity are anti-correlated.

Q23. Similarly, what two variables are most positively correlated?

Chocolate and bar are the most positively correlated.

Principal Component Analysis

Variance of PC axes:

```
pca <- prcomp(candy, scale = T)
summary(pca)
```

Importance of components:

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 |
|------------------------|---------|---------|---------|---------|---------|---------|---------|
| Standard deviation | 2.0788 | 1.1378 | 1.1092 | 1.07533 | 0.9518 | 0.81923 | 0.81530 |
| Proportion of Variance | 0.3601 | 0.1079 | 0.1025 | 0.09636 | 0.0755 | 0.05593 | 0.05539 |
| Cumulative Proportion | 0.3601 | 0.4680 | 0.5705 | 0.66688 | 0.7424 | 0.79830 | 0.85369 |
| | PC8 | PC9 | PC10 | PC11 | PC12 | | |
| Standard deviation | 0.74530 | 0.67824 | 0.62349 | 0.43974 | 0.39760 | | |
| Proportion of Variance | 0.04629 | 0.03833 | 0.03239 | 0.01611 | 0.01317 | | |
| Cumulative Proportion | 0.89998 | 0.93832 | 0.97071 | 0.98683 | 1.00000 | | |

Plotting the first two PC axes using ggplot:

```
candy_data <- cbind(candy, pca$x[,1:3])
```

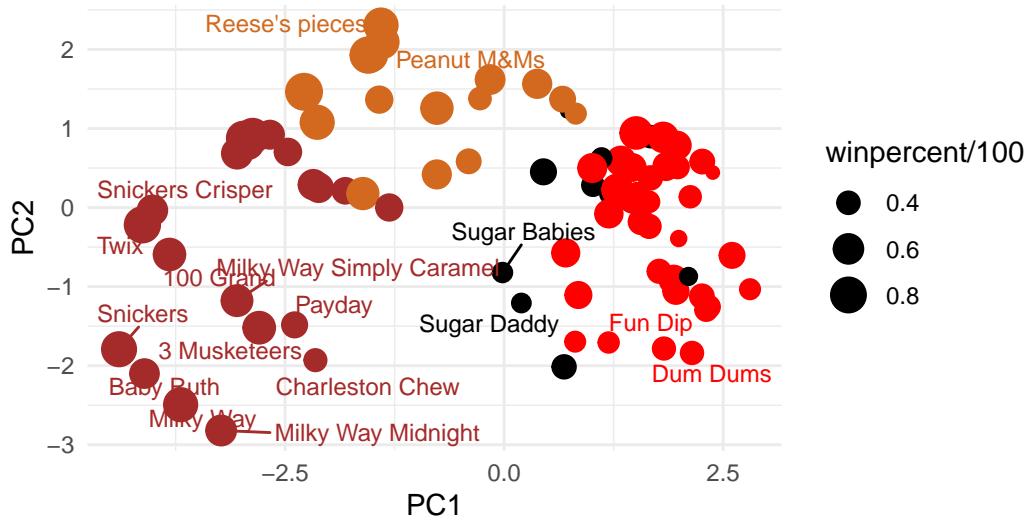
```
p <- ggplot(candy_data) +
  aes(PC1, PC2,
    size = winpercent/100,
    text = rownames(candy_data),
    label = rownames(candy_data)
  ) +
  geom_point(col = my_cols) +
  geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
    subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),
    caption="Data from 538") +
  theme_minimal()

p
```

Warning: ggrepel: 68 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown),

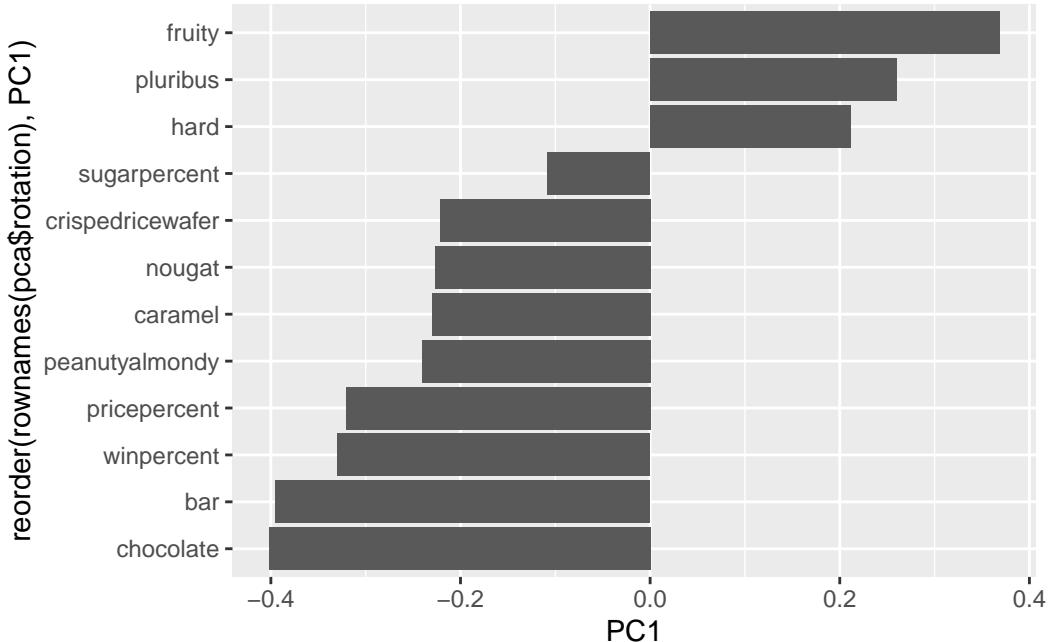


Making an interactive plot:

```
#library(plotly)
#ggplotly(p)
```

Loading scores plot for PC1:

```
ggplot(pca$rotation) +
  aes(PC1, reorder(rownames(pca$rotation), PC1)) +
  geom_col()
```



Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

Fruity and pluribus are. This is the same as the correlation plot where fruity and pluribus were strongly correlated.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

It looks like chocolates and bars are the best winning candies. This is backed up by the exploratory analysis where the bar plot shows the top winning items are chocolates or bars. This is further backed up by the correlation plot which shows that both chocolate and bars are correlated with winning and bars are correlated with chocolate. The PCA also shows that chocolate, bar, and winning are all strongly associated with the negative direction of PC1. Based on these pieces of evidence, winning candies are usually chocolates and bars.